

Programming with Bluespec SystemVerilog: Design Methodologies, Styles, Patterns and Examples

Preliminary Draft

Please do not circulate without permission from Bluespec, Inc.

Revision: 24 November 2008

Copyright © 2000 – 2008 Bluespec, Inc.

This document describes Bluespec SystemVerilog coding guidelines through a series of descriptions and examples, where the objectives are to:

- bring experienced Verilog and VHDL designers up to speed with the abilities and features of Bluespec,
- produce expected RTL with respect to interfaces, and structures,
- take advantage of the extensive libraries provided by Bluespec.

This document is organized into several sections which need not be read in sequence. Rather, each section gives hints, advice and examples on particular design or programming structures.

This document is not intended to replace the Reference Guide or the User's guide. It is intended only as a supplement to aid designers in expressing their designs in the Bluespec environment.

This document is under development.

Contents

Table of Contents	2
1 Types	4
1.1 Bit Types	4
1.2 Non Bit Types	4
1.3 Conversion Between Types	4
1.4 Types from the Bluespec Library	5
2 Designing Interfaces	8
2.1 Interface Basics	8
2.2 Sharing Signals in an Interface	9
2.3 Combining Interfaces	10
2.4 Basic Interfaces from the Bluespec Library	10
2.5 Interface Paradigms from the Library	11
3 Logic Representation	12
3.1 Sequential Element	12
3.2 Combinational Logic	12
4 System Design Examples	13
4.1 Synchronous State Machine	13
4.2 Defining Interface Methods	15
4.3 Extracting pieces from an interface	16
4.4 Using RWire to Avoid Registers and Latency	16
5 Testbenches	18
5.1 Controlling Simulation	18
5.2 Stored Test Patterns	19
5.3 Generating Random Test Patterns	19
6 Common Hints and Style Suggestions	21
6.1 Identifier Names	21
6.2 Use let Variables	21
6.3 Using types instead of ‘defines’	21
6.4 Adding Monitors	22

7	Helping the Scheduler	23
7.1	Rules versus Always Blocks	23
7.2	Alleviating Read/Write Conflicts with ConfigReg	24
7.3	Register Updates	24
8	Debugging Hints and Tips	26
8.1	Viewing Complex Data Structures	26
9	BlueSim	29
9.1	Improving Simulation Speed	29
10	Other Notes	30
10.1	Using Bluespec with other tools	30
11	Using Existing Verilog Components in a BSV Design	31
11.1	Sample Verilog Memory Import	31

1 Types

Bluespec provides a strong, static type-checking environment. Everything has a type; that is, every variable and every expression has a type. Variables must be assigned values which have compatible types. Type checking, which occurs before program elaboration or execution, ensures that object types are compatible and that needed conversion functions are valid for the context.

This section describes common types and structures as defined and used in the Bluespec environment.

1.1 Bit Types

At the lowest level, synthesizable objects can be considered as a wire or wires having some fixed width and interpretation. These correspond to Verilog `wire` and `reg`, and additionally have tighter semantics surrounding their use.

- `Bool` is a Boolean – a True or False value. The implementation is one bit, but bit-wise operations on Booleans in Bluespec are illegal. Booleans should be used for True or False values; `Bit#(1)` should be used for zero or one variables.)
- `Bit#(n)` defines a type containing `n` bits. Type `Bit` allows bit-wise operations, but `Bit#(1)` cannot be used as a `Bool`; that is, operators `&&`, `||`, and `!` are not allowed. Use `Bit#(n)` types for variables which are simple bit patterns. Note: `bit[15:0]` is a synonym for `Bit#(16)`.
- `UInt#(n)` defines an unsigned integer of `n` bits in width. Use `UInt#(n)` for variables which should be interpreted as unsigned quantities. Note: unsigned variables are always greater than or equal to zero.
- `Int#(n)` defines a signed (two's-complement) integer of width `n`. Use `Int#(n)` for variables which should be interpreted as signed quantities.

These bit type can be combined into other types through the use of user defined structures or by using the many pre-defined types in the Bluespec library. (See *Reference Manual* for details of the documentation.)

1.2 Non Bit Types

In addition to the many bit types found in Bluespec, there are many other types which are only used during compile-time evaluation, and cannot be turned into hardware bits. These types are most useful during compile time elaboration, where Bluespec allows the manipulation of objects of these types for complicated and highly parameterized designs.

- `Integer` – Integers are unbounded in size which are commonly used as loop indices for compile-time evaluation.
- `String` – Strings must be resolved at compile time.
- `Interfaces` – Interfaces can be considered a type, and hence can be passed to or returned from functions.
- `Action` – See *Reference Manual* for functions which manipulate this type.
- `ActionValue` – See *Reference Manual* for functions which manipulate this type.
- `Rules` – See *Reference Manual* for functions which manipulate this type.
- `Modules` – See *Reference Manual* for functions which manipulate this type.
- `functions` – Functions are first class objects and can be passed to other functions or modules.

1.3 Conversion Between Types

The Bluespec environment strictly checks both bit-width compatibility and type. This behavior differs from typical Verilog tools in that conversion is automatic in Verilog tools, whereas Bluespec requires explicit

conversions. To convert across types in Bluespec, the following overloaded functions should suffice. During type checking, the compiler resolves these functions to a particular type instance. Excessive type conversion usually indicates a poor choice of the basic object types used in the design.

- `pack()` – converts (packs) from various types, including `Bool`, `Int`, and `UInt` **to** `Bit#(n)`.
- `unpack()` – converts to various types, including `Bool`, `Int`, and `UInt` **from** `Bit#(n)`.
- `zeroExtend()` – increases the bit width of an expression by padding the most significant end with zeros. This function and the following two are overloaded in type `Bit`, `UInt`, and `Int`.
- `signExtend()` – increases the bit width of an expression by replicating the most significant bit.
- `truncate()` – shortens the bit width to a particular width.
- `fromInteger` – Converts from an `Integer` to any type where this function is provided in the `Literal` type-class. Integers are most often used during static elaboration since they cannot be turned into bit, hence there is no corresponding `toInteger` function.
- `valueOf` – Converts from a numeric type to an `Integer`. Numeric types are the “n’s” as used in `Bit#(n)` or the number 17 as used in `Bit#(17)`.

These conversion utilities do not incur logic overhead. The following example shows these type conversion functions in use.

```
...
    Bit#(1) aB1;
    Bool aBool = unpack( aB1 ) ;    // unpack from Bit
...
    Bit#(16) aB16 ;
    Int#(16) aI16 = unpack( aB16 ) ;    // size matters
...
    Bit#(16) bB16 ;
    bB16 = pack( aI16 ) ;    // convert back to bits
...
    // Truncates and Extends do not require any bit width information
    UInt#(16) aU16;
    UInt#(14) aU14 = truncate ( aU16 ) ;
...
    Bit#(16) aB16;
    Bit#(14) aB1 = truncate ( aB16 ) ;
...
    // Extending requires same base type
    Int#(14) aI14 ;
    Int#(16) aI16 = signExtend( aI14 ) ;
    UInt#(14) aU14 ;
    UInt#(16) aU16 = signExtend( aU14 ) ;
    Bit#(14) aB14 ;
    Bit#(16) aB16 = zeroExtend( aB14 ) ;
...
    // UInt to Int 16 conversion using pack and unpack
    UInt#(16) aU16 ;
    Int#(16) aI16 = unpack( pack ( aU16 ) ) ;
```

1.4 Types from the Bluespec Library

The Bluespec environment provides many abstractions and polymorphic data types in its library. A few common ones are detailed here, since they have high correlation to common (hardware) design paradigms.

Maybe#(a) The Maybe type encapsulates any data type a with a valid bit and further provides some functions to test the valid bit and extract the data. This is useful in places where you sometimes store data and other times the data is invalid. Some Bluespec primitives also use Maybe types in their interfaces, for example RWire.

Consider the example below which pulls data from a FIFO holding Maybe#(a_type), and then sends only valid data to the valid FIFO.

```
...

// A FIFO holding Maybe results
FIFO#(Maybe#(a_type)) result_fifo <- mkSizedFIFO (16) ;

// A FIFO for only valid results
FIFO#(a_type) valid_fifo <- mkSizedFIFO (16) ;

...
// If the result is invalid just dequeue the result
rule invalid_result (! isValid( result_fifo.first )) ;
    result_fifo.deq ;
endrule

// If the result is valid, strip the valid bit and re-queue.
rule valid_result (isValid( result_fifo.first )) ;
    valid_fifo.enq( validValue( result_fifo.first )) ;
    result_fifo.deq ;
endrule
```

Maybe types are a specialization of tagged unions (see *Reference Manual* for more details of tagged unions). Thus the tagged-union “matches” operator can be utilized to rewrite the rules in the above example.

```
// If the result is invalid just dequeue the result
rule invalid_result (result_fifo.first matches tagged Invalid) ;
    result_fifo.deq ;
endrule

// If the result is valid, strip the valid bit and re-queue.
rule valid_result ( result_fifo.first matches tagged Valid{x} ) ;
    valid_fifo.enq( x ) ;
    result_fifo.deq ;
endrule
```

Values of type Maybe can be created using either of their constructors – Valid or Invalid.

```
Results_t aResult ;
// wrap aResult into a Maybe type and enqueue it.
result_fifo.enq( Valid( aResult ) ) ;

...

// Enqueue Invalid
result_fifo.enq( Invalid ) ;
```

The recommended means to extract data from the Maybe type is to use tagged unions, or if a value or default is expected, to use the fromMaybe function. Tagged unions are shown in the example above, while the fromMaybe is show in the code below.

```
Maybe#(int) mValue ;
// Extract the value from the Maybe or else 0 if invalid.
int extr = fromMaybe( 0, mValue ) ;
```

The above code fragment is equivalent to:

```
Maybe#(int) mValue ;
int extr = 0 ;
if ( isValid( mValue ) )
    extr = validValue( mValue ) ;
```

The former is recommended.

Tuple2#(a,b) and Tuples A tuple provides an unnamed data structure typically used to hold a small number of related fields. Like other SystemVerilog constructs, tuples are polymorphic, that is, they are parameterized on the types of data they hold. Tuples are typically found as return values from functions when more than one value must be returned.

In the Bluespec environment, tuples are defined for 2 to 7 element, i.e., **Tuple2** through **Tuple7**. In most specific cases, named data structures should be used instead of tuples, since structures provide for better (self) documentation of the source code, since each field has a name in addition to its position.

```
typedef Tuple2#( Int#(32), Bit#(3) ) TestTuple ;
...

Int#(32) idata = 5;
Bit#(3) bdata = 3'b101 ;

// tuple constructor
TestTuple foo = tuple2( idata, bdata );

// extracting one field from tuple
Int#(32) idata2 = tpl_1( foo ) ;

// extracting all fields from a tuple using pattern matching
// the types for idata3 and bdata3 are determined by the context
// and checked by Bluespec's type checking system
let { idata3, bdata3 } = foo ;
```

2 Designing Interfaces

This section presents an overview of interfaces as they are used with Bluespec SystemVerilog. This section focuses on defining interfaces, and does not detail how methods can be defined. See Section 4.2 for details and examples of defining methods, since many implementations can provide the same interface.

2.1 Interface Basics

In Verilog interfaces for modules are port lists; conceptually a bundle of independent of wires which connect the module to its parent. SystemVerilog extends the port list concept by providing interfaces which encapsulate interconnect and communication thus separating communication from functionality and enabling abstraction in the design.

SystemVerilog interfaces have many similarity to Verilog modules in that interfaces must be defined and instantiated before use. Moreover, interfaces or modules can provide **tasks** and **functions** to further encapsulate communication between modules.

In the Bluespec environment, interfaces to modules are methods, that is, specific functions which may be invoked by the caller. These functions may take zero or more arguments, can return values or cause actions to occur. Bluespec interfaces provide a means to group wires into bundles with specified uses.

Bluespec classifies methods into types:

- Value Methods – These are methods which return a value to the caller. When these methods are called, there is no change of state.
- Action Methods – These are methods which cause actions (state changes) to occur. One may consider these as input methods, since they typically take data into the module.
- ActionValue Methods - These methods couple Action and Value methods, providing an action to the module which provided the method, and returning a value to the caller.

Consider the following interface definition for a FIFO-like device:

```
interface Fifo_ifc#(type any_t);
    method any_t      first_element() ;
    method Action      clear() ;
    method Action      enqueue( any_t data_in ) ;
    method ActionValue#(any_t) dequeue() ;
endinterface
```

In this example, `first_element` delivers (as an output) a value of type `any_t`, without any change to the module which provides this interface. `clear` is an Action method, which takes no argument, and returns nothing, but it does impact the state of the module – that is it causes an Action (e.g., clearing the fifo). Likewise, the `enqueue` method causes a state change without returning a value, but it does take an argument. The `dequeue` method causes a state change and returns a value.

Interfaces only specify the methods which a module provides, they do not say anything about any implementation. The above fifo-like interface can be used for fifo of any size. It can also be provided by a module which dequeues the square of the number enqueued, or one which dequeues the sum of a `n` consecutive enqueues.

Note that in this example the interface takes a type parameter; thus this interface can be used (reused) in situations which differ by data types:


```

// declaring a module with an interface of Fifo_ifc
module mkFifo1( Fifo_ifc#(any_t) ) ;
    ...
endmodule

...
// instantiate a interface with a specific type Bit#(32)
Fifo_ifc#(Bit#(32)) fifo32_inst1 <- mkFifo1 ;
...

```

Implicit Signals Note that Bluespec automatically provides appropriate handshaking signals and logic for each method of the interface. Specifically for `fifo_ifc` interface, the following handshaking signals are generated:

- `RDY_first_element` – An output which indicates that the value is valid.
- `RDY_clear` – An output which indicates that the clear method can (safely) be called.
- `EN_clear` – An input which triggers the clear method.
- `RDY_enqueue` – An output which indicates that the enqueue method can be called.
- `EN_enqueue` – An input which triggers the call to the enqueue method. Data is assumed value at `data_in`, and it is assumed (responsibility of the caller) that `RDY_enqueue` – is asserted before this signal was asserted.
- `RDY_dequeue` – An output which indicates that action value method dequeue can be called.
- `EN_dequeue` – An input signal which triggers the action in dequeue.

Bluespec provides attributes which allow designer to assert that methods are always ready (* `always_ready` *) or that the actions are always enabled (* `always_enabled` *).

2.2 Sharing Signals in an Interface

Consider an interface to a memory system as follows:

```

interface Memory_ifc ;
    // read returns data at address
    method Data_t read ( Addr_t addr ) ;
    // write datain at address
    method Action write ( Addr_t addr, Data_t datain ) ;
endinterface

```

In this definition, there would be four buses (`read_addr`, `read`(return value from read method), `write_addr`, and `write_datain`) plus three handshaking signals (`RDY_read`, `RDY_write`, and `EN_write`). It may be desirable to have the address bus shared between the two methods. This section shows some styles on how this can be done.

Using a Common Method In this style, the read and write methods are combined into one method.

```

interface Memory_ifc ;
    method Action read_write( Bool mode, Addr_t addr, Data_t datain ) ;
    method Data_t data_out() ;
endinterface

```

In this example, there are three buses (`data_out`, `read_write_addr`, and `read_write_datain`) plus three handshaking signals (`RDY_read_write`, `RDY_data_out`, and `EN_read_write`).

Using a Separate Method for the Shared Signals In this style a separate action method is added which takes in the address for both read and write methods. The implementation of the `addr` method may require some special techniques, such as described in Section 4.4 on `RWire`.

```
interface Memory_ifc ;
    method Action    addr( Addr_t addr ) ;
    method Data_t    read ( ) ;
    method Action    write ( Data_t datain ) ;
endinterface
```

In this example, there are three buses (`addr_addr`, `read`, and `write_datain`) plus five handshaking signals (`RDY_addr`, `RDY_read`, `RDY_write`, `EN_addr` and `EN_write`).

2.3 Combining Interfaces

Consider that interfaces are treated like structures (collections) of methods, then it is possible to combine separate interfaces into one, (i.e., an interface composed of other interfaces.)

```
interface Part1_ifc ;
...
endinterface
...
interface Combined_ifc ;
    interface Part1_ifc part1 ;
    interface Part1_ifc part2 ;
    interface Part2_ifc part3 ;
endinterface
```

2.4 Basic Interfaces from the Bluespec Library

The Bluespec library provides interfaces for many common hardware elements, such as Registers, FIFOs, and Register Files. These interfaces are polymorphic; that is, the interfaces are parameterized on the type of data used in the elements. Moreover, several variations of modules can provide the same interface types, thus allowing easy reuse when modules change.

Reg Interface Probably the most common interface is type `Reg#(type a)`, which is an interface to a register. The definition is polymorphic on the type of data (type `a`), and provides two methods: `_read` and `_write`. Bluespec provides a shorthand notation for reading and writing registers, because of their frequency of use. The following two action blocks are functionally equivalent in the code fragment below.

```
...
    // instantiate 2 registers
    Reg#(Data_t)    debug_datain <- mkRegU ;
    Reg#(Data_t)    datain       <- mkRegU ;

...
    // Using explicit function calls of the Reg interface
    action
        datain._write( datain._read + 500 );
        debug_datain._write( datain._read );
    endaction

    // Alternative, using Reg interface shorthand notation
    action
```

```

        datain <= datain + 500 ;
        debug_datain <= datain ;
    endaction
...

```

The Reg interface type is provided by several modules from the Bluespec Library; these include mkReg, mkRegA, mkRegU, mkConfigReg, mkConfigRegA, and mkConfigRegU.

FIFO Interface Like the Reg interface, the FIFO interface is polymorphic on its type – FIFO#(type a). It provides four methods, and its prototype is given below.

```

interface FIFO #(type a);
    method Action enq(a x1);
    method Action deq();
    method a      first();
    method Action clear();
endinterface

```

The FIFO interface type is provided by the module constructors: mkFIFO, mkFIFO1, mkSizedFIFO, and mkLFIFO. Other modules which are not FIFOs may provide this interfaces; see Section 4.3 on extracting pieces from an interface. Moreover, the FIFO type interface provides a basis for the GetPut interface, and is a common building block for many modules and interfaces.

RWires and PulseWires Although RWire is not a common hardware element, it provides a useful tool in the Bluespec environment for transporting data between methods and rules without the cycle latency of a register. The RWire interface is defined below.

```

interface RWire #(type a);
    method Action wset(a x1);
    method Maybe#(a) wget() ;
endinterface

```

The interface can be considered as a wire plus a valid signal. The wset method writes a value to the “wire” and sets the valid signal. The wget method return the value and the valid signal in a Maybe type. See Section 1.4 for a description of the Maybe and Section 4.4 for an example of using RWire.

There are two alternative interfaces for RWires, one in the common Reg interface, and the second is the PulseWire interface, which wraps a RWire without any data.

```

interface PulseWire #(type a);
    method Action send ();
    method Bool   _read() ;
endinterface

```

The use of these interfaces, and the corresponding modules constructors mkWire and mkPulseWire, are strongly encouraged over the primitive RWire.

2.5 Interface Paradigms from the Library

The Bluespec environment provides several packages, each containing interface types and associated functions which allow easy generation and connection of interfaces. See *Reference Guide* for more details.

3 Logic Representation

3.1 Sequential Element

In Verilog, a declaration of “reg” is for a variable which can be assigned in a sequential (always) block. A “reg” declaration does not imply that a hardware “Register” or a Flip-Flop will be created for this variable. It may be the case that a register is created by a synthesis tool for a variable declared as a reg, but the structure of the Verilog program determines how a variable is synthesized.

In Verilog for RTL synthesis, registers (flip-flops) are modelled and inferred by a sequential always block which has a basic structure such as:

```
reg [5:0] state, next_state ;
always @(posedge clk)
    state <= next_state ;
```

Even though both `state` and `next_state` are declared as reg, an RTL synthesis tool creates a register for `state`, and not `next_state`.

With Bluespec, declarations of registers or flip-flops as well as other state elements (such as FIFOs) are explicit by using the `mkReg` (or `mkFIFO` for FIFOs) function declarations.

```
// Create 6 bit wide register with a 0 reset value
Reg #(Bit#(6)) state <- mkReg (0) ;      // Create a register with interface state

// Create a FIFO interface and fifo
FIFO #(Bit#(6)) input_fifo <- mkFIFO ;
```

Note that in the Bluespec example, the Register and its interface are declared, but the specifics of the inputs, outputs, clock and reset port connections are not described; this task is handled by the compiler. The reading and writing of Bluespec registers, FIFOs, and other modules is accomplished through their interface methods in rules or interface methods.

3.2 Combinational Logic

In Verilog, combinational logic can be described by either continuous assignments, or with always blocks. RTL synthesis tools transform these structures into a logic function, which are then mapped to gates. In Bluespec, combinational logic is best described by functions. Moreover, Bluespec compilation has the power to automatically generate control logic for Finite State Machines (FSMs) and handshake logic between FSMs or interface methods.

Verilog continuous assignments can be thought of as functions, and combinational logic cones as a composition of functions. This conceptual model maps directly onto Bluespec, which encourages the use of functions. The significant difference between continuous assigns with and without functions is that functions can be defined once and instantiated many times, whereas continuous assigns explicitly have one instantiation in hardware.

During Verilog synthesis both combinational and sequential always blocks are transformed from its sequential execution (imperative) semantic to a functional semantic for easy synthesis into logic gates. That is, sequential if-then-else statements are transformed to muxes, with both branches evaluated in parallel. With Bluespec, the model is to divide the program into functions, using the many paradigms available in Bluespec. See Section 4 for detailed design examples.

4 System Design Examples

This section describes several design examples using Bluespec SystemVerilog.

4.1 Synchronous State Machine

With Bluespec, a state machine can be written by a series of independent rules, which the compiler will properly order and synthesize control logic. Alternately, the transitions can be combined into a function and that function instantiated in one rule, which highlights the independent and synchronous nature of the rules.

FSM with Multiple Rules

```
...

typedef enum {S0, S1, S2, S3} StateType
    deriving (Eq, Bits);

...

Reg#(StateType) curr_state <- mkReg (S0) ;

(* fire_when_enabled, no_implicit_conditions *)
rule state_S0 (curr_state == S0);
    curr_state <= S1;
endrule

(* fire_when_enabled, no_implicit_conditions *)
rule state_S1 (curr_state == S1);
    curr_state <= S2;
endrule

(* fire_when_enabled, no_implicit_conditions *)
rule state_S2 (curr_state == S2);
    curr_state <= S3;
endrule

(* fire_when_enabled, no_implicit_conditions *)
rule state_S3 (curr_state == S3);
    curr_state <= S0;
endrule

...
```

Each rule corresponds to one or more state transitions, where the rule condition is the current state and input condition for transition to occur.

Each rule is also annotated with a `fire_when_enabled` attribute. This attribute enables compile time checks to insure that other conditions (rules or method will not block the rule execution – basically a check for independence between rules.) The `no_implicit_conditions` attribute, is a compile-time assertion which insures that there are no implicit conditions (such as full or empty FIFOs) which would prevent this rule from firing. Together, these attribute check the designer’s intentions and understanding of the Bluespec model.

FSM with One Rule Alternately, the FSM may be described in one rule, which gives the full FSM behavior in a case statement.

```
Reg#(StateType) curr_state <- mkReg (S0) ;

(* fire_when_enabled *)
rule rule1_StateChange (True);
  StateType next_state;
  case (curr_state)
    S0 : next_state = S1;
    S1 : next_state = S2;
    S2 : next_state = S3;
    S3 : next_state = S0;
    default : next_state = S0;
  endcase
  curr_state <= next_state;
endrule: rule1_StateChange
```

Another One Rule FSM Example This FSM example is a minor variation of the above, where a function describes the state transitions and the rule instantiates the function. This style has the advantage that the rule bodies, which update state, and call actions are separate from the function, thus allowing easier maintenance and reuse.

```
typedef enum {S0, S1, S2, S3} StateType
  deriving (Eq, Bits);
function StateType next_state (StateType curr_state);
  let myLocalVar;
  case (curr_state)
    S0 : myLocalVar = S1;
    S1 : myLocalVar = S2;
    S2 : myLocalVar = S3;
    S3 : myLocalVar = S0;
    default : myLocalVar = S0;
  endcase
  return myLocalVar;
endfunction

...

rule rule1_StateChange (True);
  curr_state <= next_state(curr_state);
endrule: rule1_StateChange
```

Discussion The implementation style for an FSM strongly depends on scope of the FSM within the design, where the following guidelines should be used.

- If the FSM implements only the control structure, that is, it provides the next state logic, but does not detail the actions which occur within a given state, then the FSM should be implemented in one rule. This is similar to common Verilog styles, and does not take advantage of Bluespec features.
- If the FSM design includes the actions within each state, then the recommended style is one rule per transition or state arc. That is, given two transition, state A to B and state A to C, there should be a two rules, one per transition.

This style has the advantage that Bluespec can automatically handle implicit conditions of the actions, thus shortening design, implementation and verification time. For example, if some actions in the transition from A to B are not ready, then the rule does not fire, and the state transition will not occur. Bluespec collects and analyzes these action implicit condition and then builds logic for correct behavior.

The use of multiple rules per FSM is also recommend for multiple interacting FSMs, since each rule describes correct behavior allowing Bluespec to analyze and build the correct concurrent operations of the machines.

4.2 Defining Interface Methods

When defining interfaces, the prototypes for the methods are declared. However, it is up to the module which provides the interface to define the methods. This allow different modules to provide the same interfaces, yet produce very different implementations for these interfaces. This section describes some guidelines for defining methods for interfaces.

Within a module, communication between separate methods, separate rules or between rules and methods occur through module level Bluespec SystemVerilog elements, such as registers, FIFOs, or other modules. To ascertain that redundant flops are not added and latency is retained to a minimum the following guidelines should be kept in mind:

1. Combine all combinational logic from the ports to the first stage of state elements into the interface methods. e.g. `x1` and `y1` are input ports and `sum` is the first stage of state elements

```
Reg#(a)    sum <- mkReg (0) ;    // instantiate an register interface
...

// Add the two inputs storing result in register sum.
method Action calc( x1, y1 );
    action
        sum <=  x1 + y1;
    endaction
endmethod
// note: types for x1 and y1 are taken from the interface definition.
```

2. Define rules for transitions in the state elements. i.e. for designs with multiple stages of flip-flops, all the intermediate logic can be defined in terms of rules e.g. `stage1` and `stage2` are registers. Depending upon an “opcode” which is registered, `stage2` gets the value of `stage1` either as it is or inverted

```
rule calc (True)
    case(opcode)
        3'b001: begin
            stage2 <= stage1;
        end
        3'b010: begin
            stage2 <= 1 - stage1;
        end
        ...
    endcase
endrule
```

3. Good design requires that output ports are directly registered or derived only from state elements. The output or Mealy logic should be placed in the methods as well. e.g. `z1` is an output port which is high if all bits of register `product` are high

```

Reg#(Int#(16)) product <- mkRegU ;
...

method Bit#(1) z1();
    z1 = & product;           // reduction and
endmethod;
...

```

4.3 Extracting pieces from an interface

The example below creates a module which provides a FIFO interface. The module has an input-side fifo and an output-side fifo, where the output fifo is enqueued with the square of values taken from the input fifo. The module provides a FIFO interface where the enqueue methods is provided by the input-side FIFO, while the dequeue and first methods are provided by the output-side FIFO. The clear action method calls the clear methods on both fifo.

```

import FIFO::*;

// Create a module with a Fifo interface
(* synthesize *)
module squarer ( FIFO#(int) );

    FIFO#(int) inputside <- mkFIFO ;

    FIFO#(int) outputside <- mkFIFO ;

    // a rule to move the data from input to output fifo
    rule squarethis ;
        inputside.deq;
        let datain = inputside.first ;
        outputside.enq( datain * datain );
    endrule

    // methods are first-class objects and can be assigned from other methods.
    method enq = inputside.enq ;
    method deq = outputside.deq ;
    method first = outputside.first ;
    // The clear method call both fifo clear method
    method clear ;
        action
            inputside.clear ;
            outputside.clear ;
        endaction
    endmethod

endmodule

```

4.4 Using RWire to Avoid Registers and Latency

Communication between methods and rules occur through module level constructs such as registers, fifos, or other modules. When a rule needs to access data from an interface, the input (action) method can write the data to a register, and the rule can read the value in the next clock cycle. (See Section 7.2 on using a “configuration” register to avoid schedule blocking when writing registers in methods and reading them in rules.)

To avoid this extra register and the associated cycle latency, Bluespec provides the `RWire` package. Conceptually, `RWire` is a 0-delay “wire” with an appropriate interface. The interface allows one write per cycle, and the write must occur before the read.

Shared Address Bus This example shows the use of an `RWire`, via the `mkWire` module to share information between separate interface methods. This example continues the share address bus example from Section 2.2.

```
interface Memory_ifc ;
    method Action  addr( Addr_t addrin ) ;
    method Action  read ( ) ;
    method Data_t  read_value() ;
    method Action  write ( Data_t datain ) ;
endinterface

(* synthesize *)
module mkTest3 ( Memory_ifc ) ;

    // Creates a Reg interface, but it is really a zero-latency wire.
    Reg#(Addr_t) addrbus <- mkWire ;

    Reg#(Maybe#(Data_t)) data_out <- mkReg (Invalid) ;

    method Action addr( addrin ) ;
        addrbus <= addrin ;
    endmethod

    method Action read();
        $display( "calling read with address: %h", addrbus );
    endmethod

    method Action write( data );
        $display( "calling write with address: %h - data:", addrbus, data );
    endmethod

    method read_value() if ( data_out matches tagged Valid .d );
        read_value = d ;
    endmethod
endmodule
```

5 Testbenches

5.1 Controlling Simulation

One means to control a simulation is to instantiate a counter at the top level, and key actions and possibly values from the counter values. In the example below, we setup a 16 bit counter, which is then used to start the VCD dumping, rule `init`, end the simulation rule `stop`, and control some action, rule `action1`. There is one additional rule, which constantly increments the count – rule `count`.

```
// parameterized module - a common module for testing all dut
// provided they provide the same interface - type Pusher
module push_tester#(Pusher dut_ifc) ( Empty );

    // a counter for simulation
    Reg #(Bit#(16)) count <- mkReg (0) ; // make a register with a 0 reset value

    ...

    // Start dumping
    rule init (count == 0);
        $dumpvars() ;
    endrule

    // keep counting
    rule counting ;
        count <= count + 1;
    endrule

    // finish simulation
    rule stop (count > 300 );
        $finish(1) ;
    endrule

    rule action1 (count < 300) ;
        // call the actions on the interface under test
        dut_ifc.act1()
    endrule

    ...
endmodule // push_tester
```

Notice that this example also uses a module parameter of type interface `Pusher`. That is this module uses another interface as its parameter.¹

This can be invoked as shown below.

```
(* synthesize *)
module tst_fifo( Empty );

    Pusher dut1 <- mktestpush_fifo i_dut1 ;

    Empty i1 <- push_tester ( dut1 ) ;

endmodule
```

The top-level interface is “Empty”, that is, it contains only a clock and reset line in the synthesized version. For modules with “Empty” interfaces, Bluespec provides a test driver module, one which applies a reset,

¹Future version of bsc will provide a different syntax to provide interfaces to modules for the module’s use.

and then drives the clock indefinitely. The module is located in \$BLUESPECDIR/Verilog/main.v and can be invoked by vcs using the following command-line.

```
vcs $BLUESPECDIR/Verilog/main.v +define+TOP=tst_fifo +libext+.v -y $BLUESPECDIR/Verilog/ *.v
```

5.2 Stored Test Patterns

Often while writing testbenches it is necessary to store vectors or read in vectors in the testbench. This can be accomplished by declaring a RegFile which initializes its contents at start of simulation.

```
import RegFile::*;
module mkDesign();
...
  RegFile#( Bit#(5), Bit#(11) ) stimulus_io <- mkRegFileLoad ( "input_file",0,7) ;
  // Create a register file which is indexed by 5 bits, and holds 11
  // bit data.
  //Initialize the Register File from "input_file", but and only
  // create a 8 (0 to 7) by 11 bit size.

  RegFile#( Bit#(5), Bit#(11) ) regfile2 <- mkRegFileFullLoad ("input_file") ;
  // Create a register file which is indexed by 5 bits, and holds 11
  // bit data.
  //Initialize the Register File from "input_file", and create all cell
  // create a 32 by 11 bit size. (2 ^ 5 = 32)
...
endmodule : mkDesign
```

The data file input_file is the file to read in during simulation initialization, it should be in Verilog hex format.

Using Lists for Test Patterns It is not recommended to iterate over lists, to initialize objects at runtime, since computation becomes very expensive. Each assignment generates a rule and the synthesis of rules has complexity n -squared in the worst case.

5.3 Generating Random Test Patterns

This section describes a means of generating random patterns by the use of a linear feedback shift register (LFSR). LFSR can generate pseudo-random bit patterns, and are often used in self-test circuits. In this example, we use an LFSR to trigger a rule condition in a probabilistic manner.

```
...
// a LFSR for random patterns
LFSR#( Bit#(8) ) lfsr <- mkLFSR_8 i_rand ;
...
// keep counting
rule counting ;
  count <= count + 1;
  lfsr.next ;          // update the lfsr value
endrule

// action2 occurs at random time
// lfsr ranges from 1 to 256 so probability can be adjusted
rule action2 (lfsr.value() > 128 ) ;
  mypush.go( count[3:0] );
```

```
    $display( "%t - action2 %h", $time, count[3:0] ) ;  
endrule
```

Other LFSR library functions are described in “The Reference Guide”.

6 Common Hints and Style Suggestions

This section describes some common issues and restrictions when using Bluespec. Moreover, some often overlooked convenience of Bluespec are also described.

6.1 Identifier Names

A current restriction in Bluespec SystemVerilog requires that *type* names begin with an uppercase, and that variable names begin with lower case letters. Type names include typedefs, interfaces, and structures. Variable names include module names, module instance names, interface instantiations, methods names, formal and actual argument names.

If the first character of an instance name is an underscore, (`_`), the compiler will not generate the name. This is used extensively in the Bluespec libraries as it allows layers of modules inside libraries to be hidden from users of the libraries.

If a user doesn't want a module instantiation to appear in the hierarchical naming, he can start the instance name with `_`. This tells the compiler not to generate the name. For example, a module hierarchy of `foo` calling `bar` which instantiates `readme` will generate `foo_bar_readme`. If instead, the instantiation is named `_readme`, the final generated name will be `foo_bar`.

The unintended consequence may be that there are unnamed objects generated. For example, if in the BSV code the register names start with an underscore, then they will be converted into unnamed registers. To avoid this, instance names should not start with an underscore, unless the desired behavior is to *not* generate the instance name.

6.2 Use let Variables

`let` definitions are a short hand step which directs the compiler to resolve and type check variable types. `let` variable are useful to give meaningful names to intermediate values and also to share intermediate values.

```
let norm1 = normalize( bas_res ) ;
let rounded = round ( norm1 ) ;
```

as opposed to explicitly declaring the types

```
FP_I#(8,24) norm1 = normalize( bas_res ) ;
FP_I#(8,25) rounded = round ( norm1 ) ;
```

One caveat, if too many `let` declarations are used, BSC may not be able to determine some intermediate types. This is especially true when there are truncate and extension operations surrounding these declarations.

6.3 Using types instead of 'defines

The following code fragment show how typedef and variables can be used in lieu of Verilog Pre-processor commands.

```
typedef Bit#(32)    DataT;
typedef Bit#(30)    Cntr;
typedef enum { CmdAdd, CmdSub, CmdMul, CmdDiv } DOp deriving(Bits,Eq);
```

```

typedef struct {
    Cntr    counter ;
    DataT    data1;
    DataT    data2;
    DOp      opcode}  Cmd
                    deriving(Bits) ;
typedef Tuple5#(Cntr,DataT,DataT,DOp,Cntr) CmdCnt;

Integer numAdd = 5;
Integer numMul = 3;
Integer numDiv = 1;
Integer numTotal = numAdd + numMul + numDiv;

```

These definitions can be specified in a file and accessed via the `import` syntax.

6.4 Adding Monitors

To aid in design debug, there are several bsc techniques which can be used, either alone or in combination.

BSC options Use the `-keep-fires` option on the bsc command line. Verilog code generated with this option will keep all the `CAN_FIRE_<rule>` and the `WILL_FIRE_<rule>` signals, even if they can be optimized away.

Synthesis Directives By adding `(* synthesize *)` directive on modules, the Bluespec generated Verilog will have more levels of hierarchy, which may allow easier debug. Note that the scheduling phase of Bluespec may produce a different, typically less optimized, scheduler than one where the entire design is scheduled. Hence these synthesis directives should be removed before gate-level synthesis. Note that the synthesis directive can only be used on non-polymorphic modules.

Functions can also be synthesized as a Verilog module, by preceding the function with a `(* nonline *)` attribute. Like the `(* synthesize *)` directive the generated Verilog may be less efficient when using this attribute, since optimizations cannot take place across these boundaries. Parameterized functions cannot be synthesized directly into Verilog.

Adding Hardware Monitors Yet another technique is to add a “monitor signal” or Probe. Probes are not optimized nor removed in Bluespec, so the signal will remain for debug. For example,

```

...
Probe#(Data_t)  debug_datain1 <- mkProbe ;
...
rule someRule ( ... ) ;
    ...
    debug_datain1 <= ...
    ...
endrule
...

```

Note that these types of monitors can only be added in a module context, that is not in functions, and can only be assigned in action contexts, such as rules, or action methods.

7 Helping the Scheduler

7.1 Rules versus Always Blocks

Consider the following Verilog code fragment, and a corresponding (incorrect) Bluespec model.

```
reg [31:0] burst_count, dest_burst_size ;
reg      burst_in_progress ;
...
always @(posedge clk)    // burst_in_progress_stop
begin
    if (burst_in_progress && (burst_count==dest_burst_size))
        burst_in_progress <= False;
    ...
end

always @(posedge clk)    // burst_count;
begin
    burst_count <= burst_in_progress ? burst_count+1 : 0;
    ...
end
```

Incorrect Bluespec model:

```
...
Reg#(Data_t) dest_burst_size <- mkReg (32'h0) ;
Reg#(Data_t) burst_count    <- mkReg (0) ;
Reg#(Bool) burst_in_progress <- mkReg (False) ;

rule burst_in_progress_stop (burst_in_progress && (burst_count==dest_burst_size));
    burst_in_progress <= False;
    ...
endrule

rule burst_counting ;
    burst_count <= burst_in_progress ? burst_count + 1 : 0;
    ...
endrule
```

At first inspection the Bluespec code may look equivalent, but the scheduler sees the cross correlation between the rules. That is, Rule `burst_counting` reads `burst_in_progress` and updates register `burst_count` Rule `burst_in_progress` reads `burst_count` and sets `burst_in_progress`.

For the scheduler to allow the rules to happen in parallel, they must be able to occur in any order and still give the same result. If Rule `burst_counting` happens first, then register `burst_count` can be updated such that rule `burst_counting` become disabled; a similar argument happens in the other directions.

A corrected model is below, which combines the two rules into one rule which only occurs when `burst_in_progress` is True. Note that the two actions in the rule can be listed in either order, since all actions in one rule occur concurrently.

```
...
Reg#(Data_t) dest_burst_size <- mkReg (32'h0) ;
Reg#(Data_t) burst_count    <- mkReg (0) ;
Reg#(Bool) burst_in_progress <- mkReg (False) ;

rule burst_in_progress_stop (burst_in_progress) ;
    burst_in_progress <= burst_count != dest_burst_size ;
    burst_count <= (burst_count != dest_burst_size) ? burst_count+1 : 0;
```

```
...
endrule
```

7.2 Alleviating Read/Write Conflicts with ConfigReg

According to term-rewriting semantics, rules are thought of as firing one-at-a-time sequentially. From this point of view, the scheduling logic’s job is not only to generate this sequence, but also to divide it into ”chunks”, so all the rules in a chunk fire during the same clock. However, if the rules involve several registers, all the reads from these registers will be performed near the beginning of the clock cycle, and all the writes at the end. It follows that any rule (or method) which reads a register must occur earlier in the TRS sequence as one which writes to it. Sometimes, however, there are conflicting constraints between two rules; if so, those two rules cannot be scheduled in the same chunk (that is, fired on the same clock).

An example of this occurs when a method writes to a register which is read by an internal rule. For technical reasons a module’s methods are always scheduled first in a chunk, earlier than all the internal rules, and with greater urgency. So this combination of circumstances would mean that the method and the rule can never fire together, and the rule will be inhibited whenever the method fires.

This situation often occurs with configuration registers, which are written by an external method (connected to some local bus mechanism) which writes the register on every clock (though the value written changes very rarely). In this case a rule reading the register would be starved out completely, and never fire.

The ConfigReg package was designed to address this issue. A ConfigReg has the same interface as a Reg, and the Verilog implementation for mkReg and mkConfigReg are also identical. The only difference is the scheduling information specified in the ”import ”BVI” wrapper – for ordinary registers, reads must occur before writes (as described above); for ConfigReg’s, the two methods are conflict-free, so reads may occur before or afterwards.

Thus the conflict described above no longer occurs. The downside is that the value read, whether the read occurs before or after a write within the same chunk, will always be the contents of the register at the beginning of the chunk – that is, the value might be slightly stale. In the case of configuration registers this does not normally matter – the parameters are set very rarely, and a one-clock difference is neither here nor there.

7.3 Register Updates

In Bluespec registers are considered atomic units with regard to their updates. That is, if any bit of a register is updated, then the entire register is considered updated. As an example, consider the follow code fragment. There are two rules which each update a separate byte of a 32 bit register. Bluespec semantics say that these two rules conflict, since both are updating the same register.

```
Reg#(Bit#(32)) reg1 <- mkReg(0) ;

rule updateByte1 ( r1 ) ;
  reg1[7:0] <= 8'd4 ;
endrule

rule updateByte3 ( r2 ) ;
  reg1[23:16] <= 8'd5 ;
endrule
```

A better methodology is to define four 8-bit registers, and update the bytes a needed, if this is the intended behavior. This style of coding has been observed to have better synthesis results.


```

Reg#(Bit#(8)) regA1 <- mkReg(0) ;
Reg#(Bit#(8)) regA2 <- mkReg(0) ;
Reg#(Bit#(8)) regA3 <- mkReg(0) ;
Reg#(Bit#(8)) regA4 <- mkReg(0) ;

rule updateByte1A ( r1 ) ;
  regA1 <= 8'd4 ;
endrule

rule updateByte3A ( r2 ) ;
  regA3 <= 8'd5 ;
endrule

```

Convenience functions can also be written to make the reading and writing of these registers easier.

```

function Action updateRegA( Bit#(32) din ) ;
  return
  action
    regA1 <= din[7:0] ;
    regA2 <= din[15:8] ;
    regA3 <= din[23:16] ;
    regA4 <= din[31:24] ;
  endaction;
endfunction

let regA = { regA4, regA3, regA2, regA1 } ;
...
rule updateAllRegA ( r5 ) ;
  updateRegA( regA + 1 ) ;
endrule

```

8 Debugging Hints and Tips

8.1 Viewing Complex Data Structures

Advanced data structures, such as vectors, lists, and structures, are difficult to view during simulation and debugging since they are converted into bits during compilation. You can use the `Probe` primitive which is provided in the Bluespec library, to ensure that signals of interest are not optimized away by the compiler and are given a known name. The `Probe` primitive is used just like a register, except that only a `write` method exists. Since `reads` are not available, a `Probe` does not impact scheduling. associated signal will be named just like the port of any Verilog module, in this case `<instance_name>$PROBE`, but no actual `Probe` instance will be created.

The first step is to define a specialized `probe` module based on the module `mkProbe` provided in the BSV library. In this example we'll call our module `mkS_Probe`. This module instantiates a `mkProbe` module for each field of the structure and defines the `write` method to include each field of the structure.

The specialized module, `mkS_Probe`, has to be kept in sync with the structure definition. For each field added to the structure, two lines must be added to the `mkS_Probe` module definition; one for the instantiation and one for the write. A unique module must be written for each unique data structure, but they all use the same basic techniques.

```
// Probe library is needed for this technique
import Probe :: * ;

// Define a struct, for which we want to see the internals
typedef struct {
  Bit#(sx)   xx;
  Bit#(st)   tt;
  Bool      zz ;
} MyS #( type sx, type st )
deriving ( Bits, Eq );

// Here we define a module specific to the MyS structure above.
// This module provides the Probe interface, and hence works exactly like the
// mkProbe module.
// This module should be keep in sync with the structure definition
// above. 2 lines are need for each new field added - one to
// instantiate the mkProbe module and one for the write.

module mkS_Probe( Probe#(MyS#(x,y)) ) ;

  // We need to instantiate a mkProbe module for each field of the
  // structure which we wish to probe. By using the let "type", we
  // leave the type determination to the bsc compiler
  let xx <- mkProbe ;
  let tt <- mkProbe ;           // new fields added in structure above
  let zz <- mkProbe ;           // must be added here

  // The Probe interface has one method, which we must provide in
  // this module. A _write action causes all the sub-probes to be
  // written as well.
  method Action _write( s ) ;
    xx <= s.xx ;
    tt <= s.tt ;           // new fields added in structure above
    zz <= s.zz ;           // must be added here
  endmethod
```

```
endmodule
```

Once we've defined the `mkS_Probe` module, we can use it to view the signals in the structure. The remainder of this code shows an example of using `mkProbe` and `mkS_Probe`.

```
// Define an interface and module for the example.
interface Foo ;
  method Bit#(4) mx;
  method Bit#(7) mt ;
  method MyS#(4,7) ms ;
endinterface

(* synthesizable *)
module mkTest( Foo ) ;

  let init_s = MyS{ xx: 0, tt:1, zz: True } ;

  // Create a register to hold the structure listed above.
  Reg#(MyS#(4,7)) r <- mkReg( init_s ) ;

  // Define a regular probe and the specialized probe modules
  let rprobe <- mkProbe ;
  let rprobe_current <- mkS_Probe ;
  let rprobe_new <- mkS_Probe ;

  // Probe can only be written in an action context, - e.g. a rule
  // or an action method.
  rule r1 ( True );
    rprobe <= r ;                // Write the regular probe
    rprobe_current <= r ;        // Write the specialized probe with
                                // the current value

    let ns = MyS{xx: r.xx + 5,
                  tt: r.tt + 3,
                  zz: r.zz } ;
    rprobe_new <= ns ;           // Write the specialized probe
                                // with the "new" value

    r <= ns ;
  endrule

  method mx = r.xx ;
  method mt = r.tt ;
  method ms = r ;
endmodule
```

After compiling, the generated Verilog will contain signals named `<instance_name>$PROBE`, however no actual `Probe` instance will be created. The only side effects of a BSV `Probe` instantiation relate to the naming and retention of the associated signal in the generated Verilog.

In the generated Verilog code, the following code segment will be found.

```
// probes
assign rprobe$PROBE = r ;

assign rprobe_current_xx$PROBE = r[11:8] ;
assign rprobe_current_tt$PROBE = r[7:1] ;
assign rprobe_current_zz$PROBE = r[0] ;

assign rprobe_new_xx$PROBE = x__h365 ;
assign rprobe_new_tt$PROBE = x__h385 ;
```

```
assign rprobe_new_zz$PROBE = r[0] ;
```

The first line corresponds to the `rprobe`, and the is assigned the current value of `r[11:0]`. The structure information is not present. The next 3 lines show the `rprobe_current` instance of `mkS_Probe` module, while the last 3 lines show the `rprobe_new` instance. Note that within each of these instances the specific structure fields are presented, which will aid in the observation of the design. For net-list synthesis, these signals will be removed, since they are not connected to further down-stream logic.

9 BlueSim

This section focuses on BlueSim, the native Bluespec simulator.

9.1 Improving Simulation Speed

The user has a great deal of control over how well the simulation execution performs. A poorly written model or testbench can slow down the simulation by orders of magnitude.

Initialization

Use the **RegFile** primitive to create and initialize data arrays with very low overhead.

Avoid any initialization which requires runtime actions or rules. These rules must be checked at every clock in the simulation.

Avoid creating data at initialization time for the same reason.

Rule Writing

Any source code which causes many (hundreds or thousands) of rules to be created will execute slowly.

Avoid using rules whose predicate is **True**. These will be scheduled on every tick. If possible, have all rules depend on computed conditions.

One large factor in simulation overhead is scheduling the rule firings. Using simple predicates and a judicious economy of rules will greatly speed simulation.

Computations

Long case statements generate complex code. If possible, use a **RegFile** or a computation to derive a code, then use the code to control the simulation behavior.

The compiler is very good at finding common subexpressions and other normal optimizations, so the user need not worry about these.

10 Other Notes

10.1 Using Bluespec with other tools

When using a simulator and other Design automation tools, Verilog modules from the Bluespec library will need to be included during elaboration. The examples below show how this can be accomplished via command line or command script

Simulators and other tool Simulators typically require that the Verilog library directory is specified. Bluespec provides Verilog definitions for these “primitive” modules in `$BLUESPECDIR/Verilog/`, which be accessed in most simulators as:

```
vcs,ncverilog +libext+.v -y $BLUESPECDIR/Verilog/ *.v
```

Design Compiler can access Bluespec’s Verilog library by adding the following to command scripts or in the `/.synopsys_dc.setup` file.

```
bluespecdir = get_unix_variable("BLUESPECDIR")
search_path = search_path + {bluespecdir + "/Verilog"}
```

Be sure to ungroup the library elements, especially registers, since DC can simplify the register instantiate if constant inputs are known `compile -ungroup_all`

May want to set DC variable’s `compile_sequential_area_recovery` to true, since this instructs dc to reconsider its choice of sequential cells.

11 Using Existing Verilog Components in a BSV Design

The `importBVI` statement is used to include existing Verilog components in a BSV module allowing the designer to reuse existing components and design components for reuse. The *wrapper*, or code around the Verilog program, can be polymorphic to allow reuse in a variety of designs. Since a polymorphic module cannot be synthesized, for synthesis, a non-polymorphic instance of the module is instantiated in which values for the parameters are specified.

11.1 Sample Verilog Memory Import

Presented here is a simple example importing a Verilog program, in this case a Verilog SRAM, into BSV. A polymorphic SRAM wrapper, `mkVerilog_SRAM_model` is defined, providing the polymorphic interface `SRAM_Ifc`. For synthesis, a non-polymorphic instance of this module, `mkSRAM_5_32` is instantiated.

First the interface which defines the Verilog SRAM as seen by the BSV clients is defined.

```
package SRAM_wrapper;

// Interface to the Verilog SRAM, as seen by BSV clients

interface SRAM_Ifc #( type addr_t, type data_t );
    method Action      request ( addr_t address, data_t data, Bool write_not_read );
    method data_t      read_response;
endinterface: SRAM_Ifc
```

The above interface is polymorphic, allowing the instantiation to define the address and data width variables. By defining the wrapper as polymorphic, the design can be reused with different size SRAMS.

Next, using the `importBVI` syntax (for detailed syntax see the Reference Manual), import the Verilog module. It provides the `SRAM_Ifc` interface defined above.

```
import "BVI" mkVerilog_SRAM_model =
    module mkPolymorphic_SRAM #(String filename) (SRAM_Ifc #(addr_t, data_t))
        provisos(Bits#(addr_t, addr_width),
                  Bits#(data_t, data_width));

        // Parameters passed to the Verilog model
        parameter FILENAME      = filename;
        parameter ADDRESS_WIDTH = valueOf(addr_width);
        parameter DATA_WIDTH   = valueOf(data_width);

        // Default clk and no reset passed to the Verilog model
        default_clock clk(clk);
        no_reset;

        // Verilog ports corresponding to method args, results and controls
        method      request (v_in_address, v_in_data, v_in_write_not_read)
                      enable (v_in_enable);
        method v_out_data read_response;

        // Both methods can be called in the same clock,
```

```

    // but read_response must logically precede request
    schedule (read_response) SB (request);
endmodule

```

The name `mkVerilog_SRAM` is the optional name of the Verilog module to be imported; the `.v` file name. If it is not provided it is assumed to be the same as the BSV name of the module (in this example `mkPolymorphic_SRAM`). Because the two names are not the same, the name of the Verilog module is required.

The parameter `FILENAME` refers to the file specifying the initial contents of the memory.

The above SRAM module is polymorphic, so it is not synthesizable. We can make it synthesizable by providing specific values to the address and data width parameters. You can make similar SRAM instances for other widths, following this pattern.

The module `mkSRAM_5_32` below defines a non-polymorphic instance of the previous module, by specifying particular values for the parameters. This also specifies a particular file name, `mem_init.data`, for the initial contents of memory.

```

typedef Bit#(5)  A5;
typedef Bit#(32) D32;

(* synthesize, always_ready = "read, write", always_enabled = "read" *)

module mkSRAM_5_32 (SRAM_Ifc #(A5, D32));
    // Instantiate the polymorphic SRAM wrapper
    // the initial file name is specified as the parameter
    SRAM_Ifc #(A5, D32) sram <- mkPolymorphic_SRAM ("mem_init.data");

    return sram;
endmodule: mkSRAM_5_32

endpackage: SRAM_wrapper

```

Below is the generated Verilog RTL for the module `mkSRAM_5_32`.

```

//
// Generated by Bluespec Compiler, version 3.8.69 (build 8868, 2006-08-07)
//
// On Thu Oct  5 08:34:34 EDT 2006
//
// Method conflict info:
// Method: request
// Sequenced after: read_response
// Conflicts: request
//
// Method: read_response
// Conflict-free: read_response
// Sequenced before: request
//
//
// Ports:
// Name                                I/O  size  props

```



```

// RDY_request          0      1 const
// read_response        0     32
// RDY_read_response    0      1 const
// CLK                  I      1
// RST_N                I      1 unused
// request_address      I      5
// request_data         I     32
// request_write_not_read I      1
// EN_request           I      1
//
// No combinational paths from inputs to outputs
//
//

`ifdef BSV_ASSIGNMENT_DELAY
`else
`define BSV_ASSIGNMENT_DELAY
`endif

module mkSRAM_5_32(CLK,
    RST_N,

    request_address,
    request_data,
    request_write_not_read,
    EN_request,
    RDY_request,

    read_response,
    RDY_read_response);
input  CLK;
input  RST_N;

// action method request
input  [4 : 0] request_address;
input  [31 : 0] request_data;
input  request_write_not_read;
input  EN_request;
output RDY_request;

// value method read_response
output [31 : 0] read_response;
output RDY_read_response;

// signals for module outputs
wire [31 : 0] read_response;
wire RDY_read_response, RDY_request;

// submodule sram
wire [31 : 0] sram$in_data, sram$in_out_data;
wire [4 : 0] sram$in_address;
wire sram$in_enable, sram$in_write_not_read;
mkVerilog_SRAM_model #(.FILENAME("mem_init.data"),
    .ADDRESS_WIDTH(5),

```

```

.DATA_WIDTH(32)) sram(.clk(CLK),
    .v_in_address(sram$v_in_address),
    .v_in_data(sram$v_in_data),
    .v_in_write_not_read(sram$v_in_write_not_read),
    .v_in_enable(sram$v_in_enable),
    .v_out_data(sram$v_out_data));

// action method request
assign RDY_request = 1'd1 ;

// value method read_response
assign read_response = sram$v_out_data ;
assign RDY_read_response = 1'd1 ;

// submodule sram
assign sram$v_in_address = request_address ;
assign sram$v_in_data = request_data ;
assign sram$v_in_enable = EN_request ;
assign sram$v_in_write_not_read = request_write_not_read ;
endmodule // mkSRAM_5_32

```